

Introduction to Artificial Intelligence (CS470): Assignment 3

Deadline:

Tuesday 2nd May, 2023

Setup

Google Colaboratory

Please download a starter code containing Colab notebooks [code address](#).

1. Unzip the starter code `.zip` file. You should see a `CS470_IAI_2023_Spring` folder. Create a folder in your personal Google Drive and upload `CS470_IAI_2023_Spring` folder to the Drive folder.
2. Each Colab notebook (e.g., files ending in `.ipynb`) corresponds to an assignment problem. In Google Drive, double-click on the notebook and select the option to open with Colab.
3. Once you have completed the assignment problem, you can save your edited files back to your Drive and move on to the next problem. Please ensure you are periodically saving your notebook File → Save so that you don't lose your progress.

1 Markov Decision Process [50 pts]

In this section, you design a Markov decision process (MDP) for a toy environment, called *grid-world*, which is often used for reinforcement learning (RL). The environment is a stochastic version of the pre-built discrete *gridworld* environment from *OpenAI Gym*. In order to represent a task in the environment, you define an MDP, $\langle \mathcal{S}, \mathcal{A}, \mathcal{T}, r, \gamma \rangle$, where \mathcal{S} , \mathcal{A} , \mathcal{T} , r , and γ are states, actions, transition probabilities, a reward function, and a discount factor, respectively. We particularly use a 8×10 size of *gridworld* environment, where the coordinate of left-top and right-bottom cells are $[0, 0]$ and $[7, 9]$, respectively. An agent can move onto one of the four nearest cells or stay. Please, fill your code in the blank section following the “PLACE YOUR CODE HERE” comments in the `CS470_Assignment3.problem.ipynb` file following the subproblems below.

TRANSITION MODEL: Implement a stochastic transition model of the environment dynamics. You need to fill out the `transition_model()` function, which returns a list of transition probabilities over the next states given a state and an action. In order to define the transition model, you have to consider following rules:

- The agent has five possible movements: stay, up, down, left, and right (see Fig.1),
- The agent is not allowed to move off the grid or move on an obstacle (grey); If the agent tries to move off the grid or move on an obstacle, it will end up staying at the previous cell,
- The agent applies a selected action with a probability $1 - \epsilon$,
- The agent (uniform) randomly applies an action among unselected actions with a (noise) probability $\epsilon/4$,
- The agent terminates the episode when it reaches goal or trap (red) cells, and
- The agent cannot leave goal or obstacle cells.

Note that we use predefined $\epsilon = 0.05$ on the assignment IPython notebook.

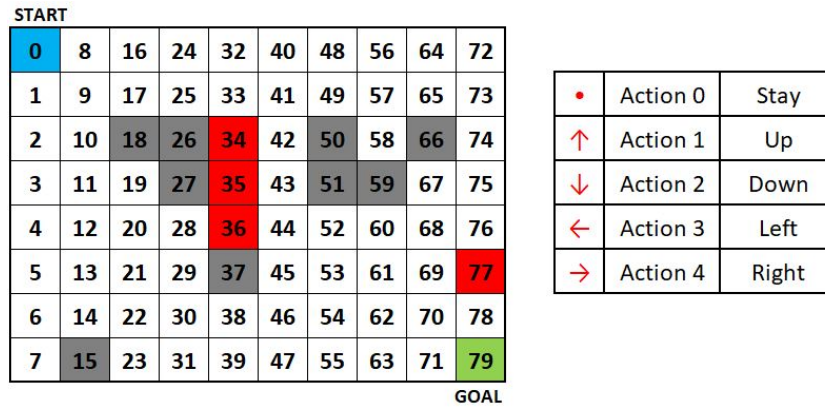


Figure 1: An exemplar *gridworld* environment with available actions. An agent cannot move on an obstacle (grey). An episode ends when the agent reaches either a trap (red) cell or a goal (green) cell. In this environment, the cell with index 1 and 78 corresponds to $[1, 0]$ and $[6, 9]$, respectively.

REWARD FUNCTION: Design a reward function by filling out the `compute_reward(s,a,s')` function, where the arguments represent a current state, an action, and a next state. The function returns a positive reward $+5$ when the agent reaches a goal (terminal state in the code), a negative reward -10 when the agent reaches a trap, otherwise 0. In addition, the function provides a step penalty of -0.1 given any action.

STEP FUNCTION: Implement the `step()` function that takes an action and applies it to the environment. The applied action leads to a stochastic transition to the next state. In details, the `step` function performs followings:

- sample a next state according to the transition function,
- calculate a reward value,
- update the current state to the next state, and
- return a result tuple, (next state, reward, termination signal, information).

TERMINATION: We terminate an episode when an agent is in either goal or trap cells. To deal with the episode termination conditions, you have to implement the `is_done()` function that returns a Boolean termination signal that is whether the current episode has to be terminated or not.

Finally, you are ready to make your agent interact with the environment for RL. On your report,

1. [15 pts] print out the transition probabilities to all the next states given
 - a current state $[3, 1]$ and a selected action “Down”,
 - a current state $[0, 6]$ and a selected action “Right”,
 - a current state $[3, 5]$ and a selected action “Right”,
2. [5 pts] plot the resulting histogram of returns produced by a dummy policy in the IPython notebook for 100 episodes,
3. [5 pts] plot the distribution of trajectories produced by a dummy policy in the IPython notebook for 300 episodes,
4. [25 pts] attach your implemented code from `transition_model()`, `compute_reward()`, `step()`, and `is_done()` functions.

2 Dynamic Programming [50 pts]

In this problem, you implement and analyze a representative DP algorithm: value iteration (VI).

2.1 Value Iteration (VI) [30 pts]

Implement the VI algorithm for the stochastic *gridworld* environment by filling out the `ValueIteration` class. The VI algorithm iteratively and directly updates each state value $V_t(s)$ at a time step t given a transition model \mathcal{T} :

$$V_{t+1}(s) = \max_a \mathbb{E}_{s' \sim \mathcal{T}(s,a)} [r(s, a, s') + \gamma V_t(s')]. \quad (1)$$

VI stops to update the values when the maximum update error Δ is lower than a certain threshold θ , where $\Delta \leftarrow \max(\Delta, \|V_{t+1}(s) - V_t(s)\|) \quad \forall s \in \mathcal{S}$. (See details on the RL book p.83.) On your report, please

- write down the state values of the first 8 states of the *gridworld* environment¹,
- overlay the best action at each state based on the state-action values,
- plot the distribution of trajectories produced by the trained policy for 100 episodes, and
- attach your implemented code from `value_iteration()` and `get_action()` functions on your report.

2.2 Comparison under different transition models [20 pts]

Suppose that the probability of taking a random action is 10% (i.e., $\epsilon = 0.1$). This transition model will affect the **exploration** process of VI. Thus, in this part, you are asked to compare/analyze the effects when you run VI with $\epsilon = 0.1$ and $\epsilon = 0.4$. On your report, please

- plot the expected returns per ϵ value with respect to the number of iterations until convergence (two graphs or one unified graph),
- overlaying the best actions at each state per ϵ value (two visualizations),
- plot the distribution of trajectories produced by the trained policy for 100 episodes per ϵ value (two visualizations), and
- compare/analyze the effect of different ϵ based on the above results.

Note that analysis does not mean you have to write down a long report. Scientific writing requires a concise delivery of your thoughts/results.

Submission Guide

2.3 Submission Requirement

You have to submit two types of materials: report and code.

- Report: Write a report answering all the questions in the assignments as a **PDF** file.
- Code: Change your code file name to `cs470_yourname_studentID_idx.ipynb` in Colab, where `idx` is the index of the problem. For example, after modifying the file `RL_problem_1.ipynb`, save as `cs470_yourname_studentID_1.ipynb`, and for the `RL_problem.ipynb`, save as `cs470_yourname_studentID.ipynb`. Download and save on your machine. Please make sure that the submitted notebooks have been saved and the cell outputs are visible.

Generate a zip file of your code and report, then save your zip file as `cs470_yourname_studentID.zip`. Please submit the `.zip` file via KLMS.

2.4 Academic Integrity Policy

This is homework for each student to do individually. Discussions with other students are encouraged, but you should write your own code and answers. Collaboration on code development is prohibited. There will be given no points in the following cases:

¹The first 8 states are the top row of states: s_0, s_1, \dots, s_7

- Plagiarism detection
- Peer cheating
- The incompleteness of the code
- The code does not work